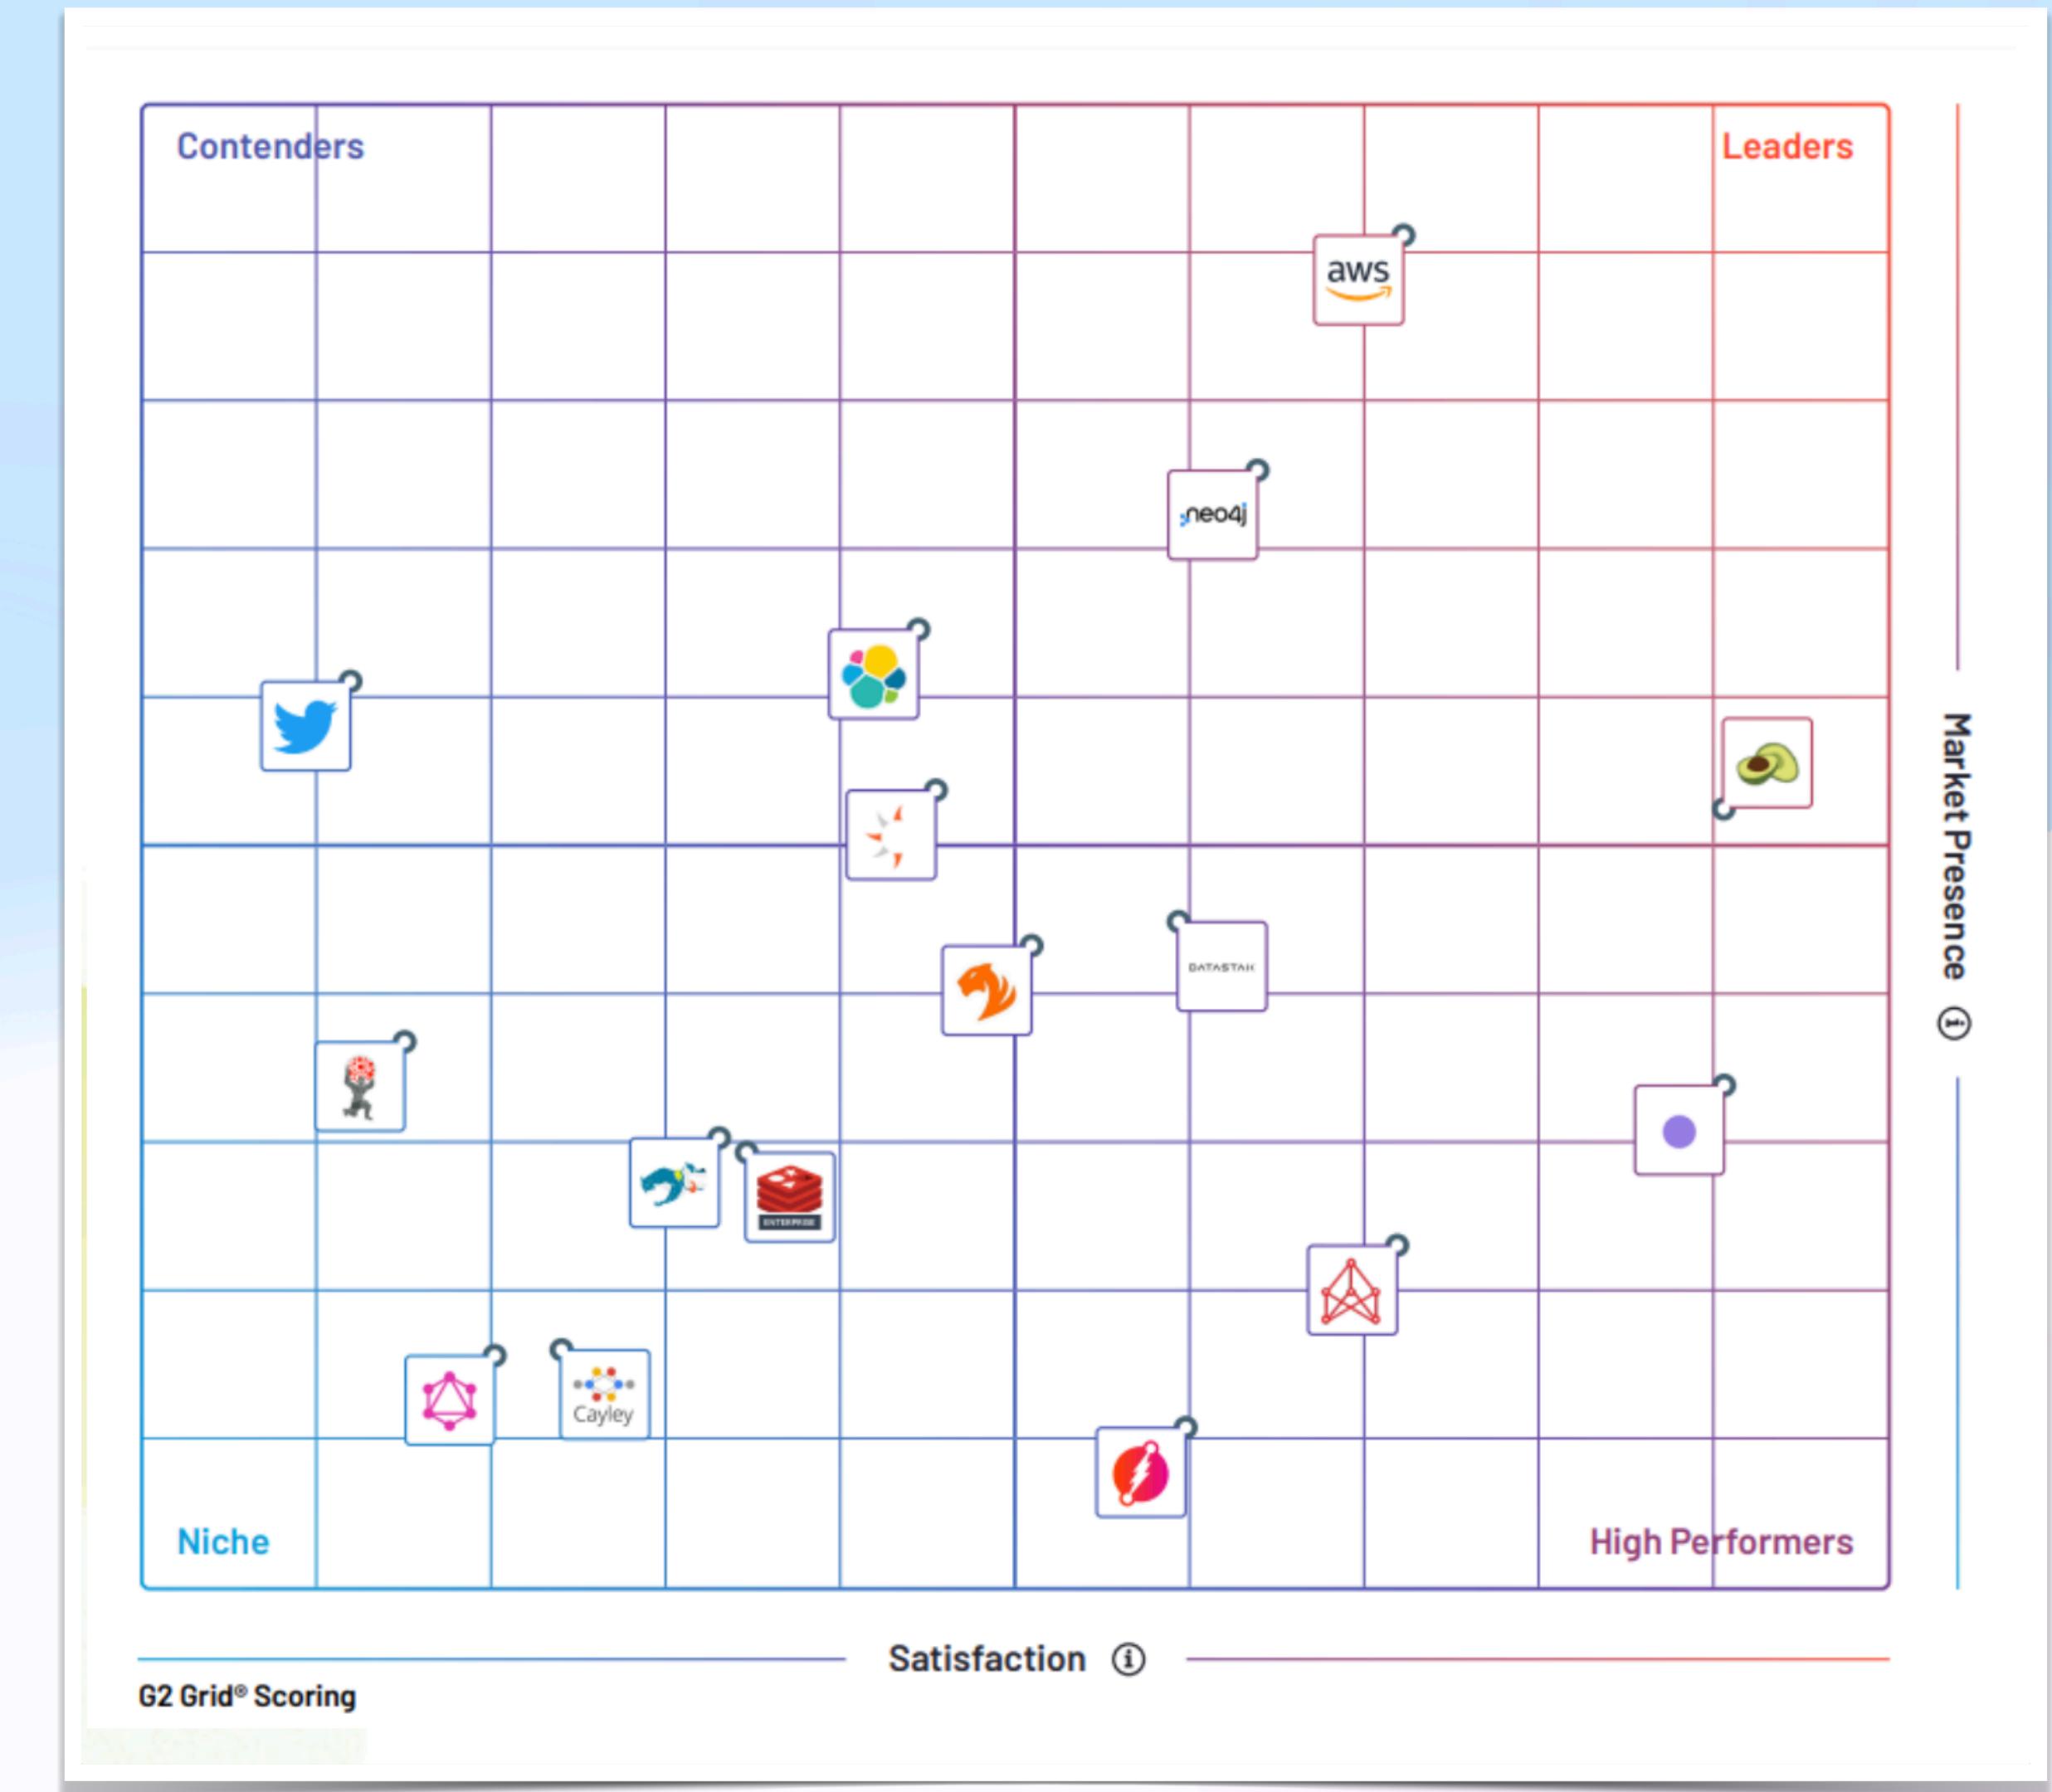


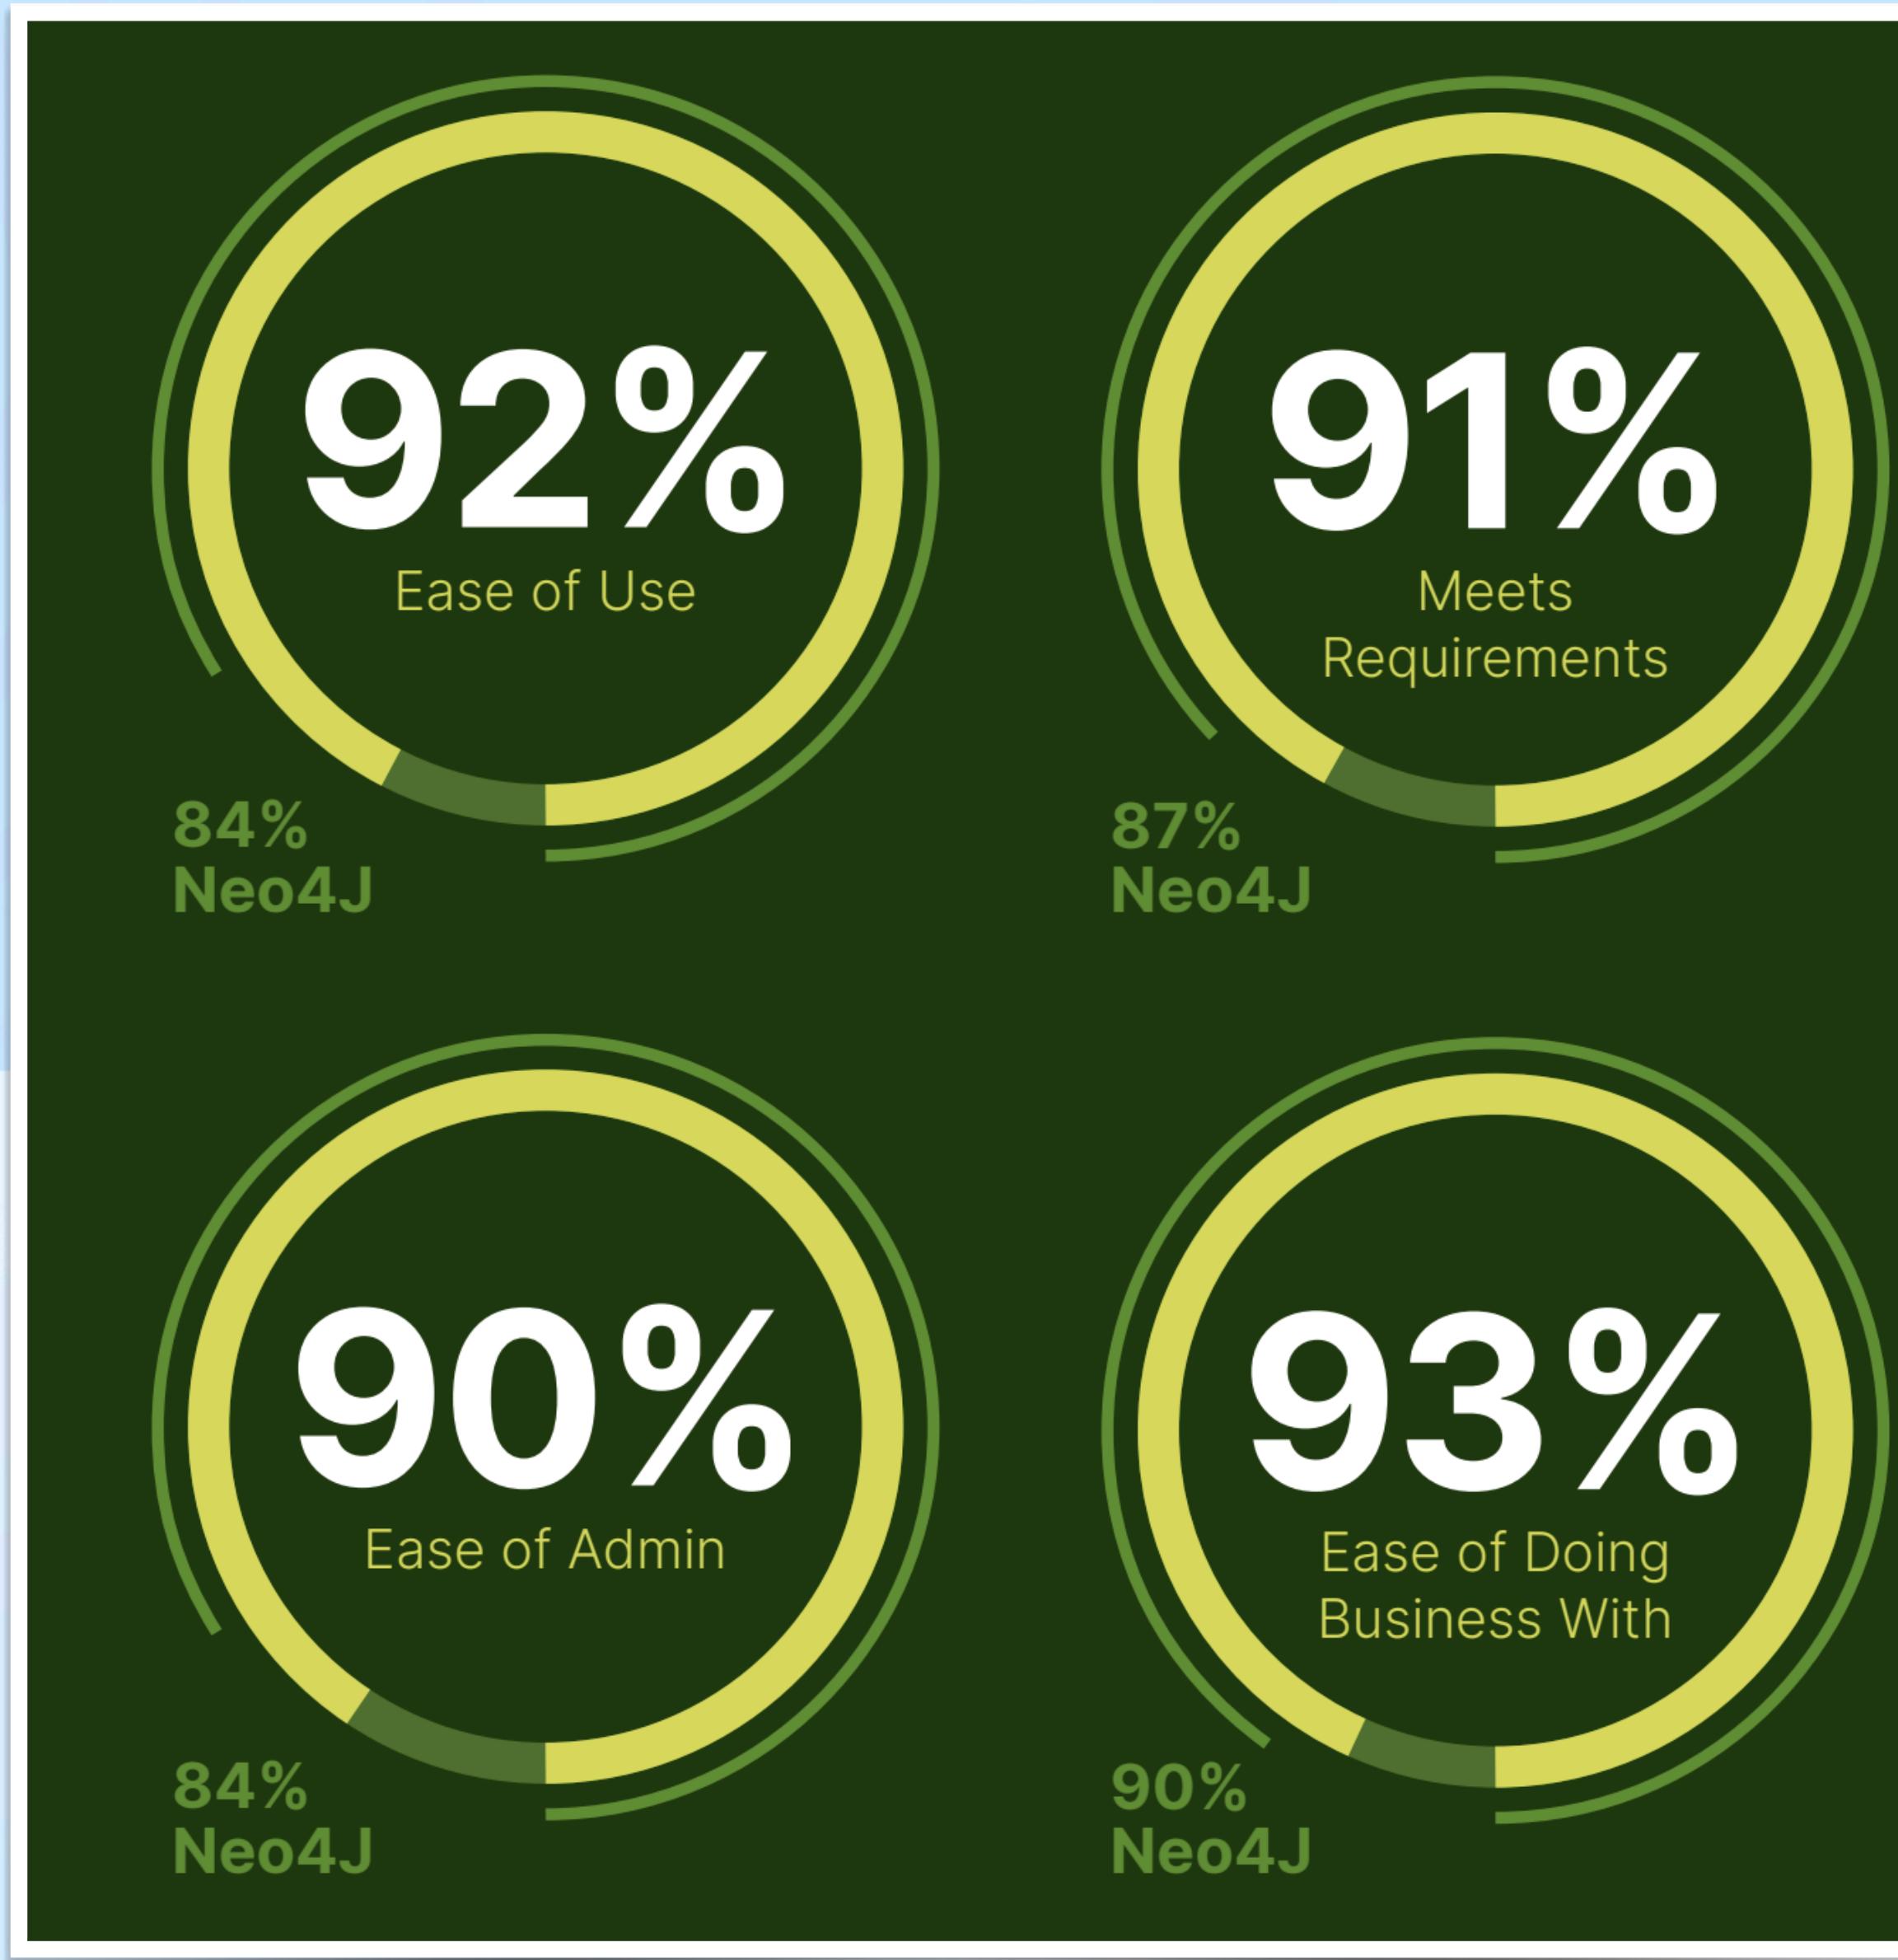
ArangoDB

final milestone struggles and victories

Mikołaj Wielgos, 24 Jan 2025

Rated the #1 Graph Database in the world for Winter 2025





"Beautiful query language,
best performance"

G2 Reviewer



"**Great choice** for multi-model
connected data"

G2 Reviewer



"ArangoDB is the **modern solution**
to the modern problem!"

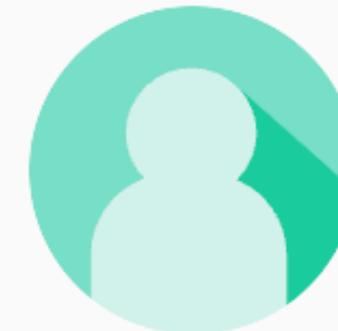
G2 Reviewer





"Great DB for
real-world data"

G2 Reviewer



Verified User in Hospital & Health Care ⓘ

Enterprise (> 1000 emp.)

Validated Reviewer ✓

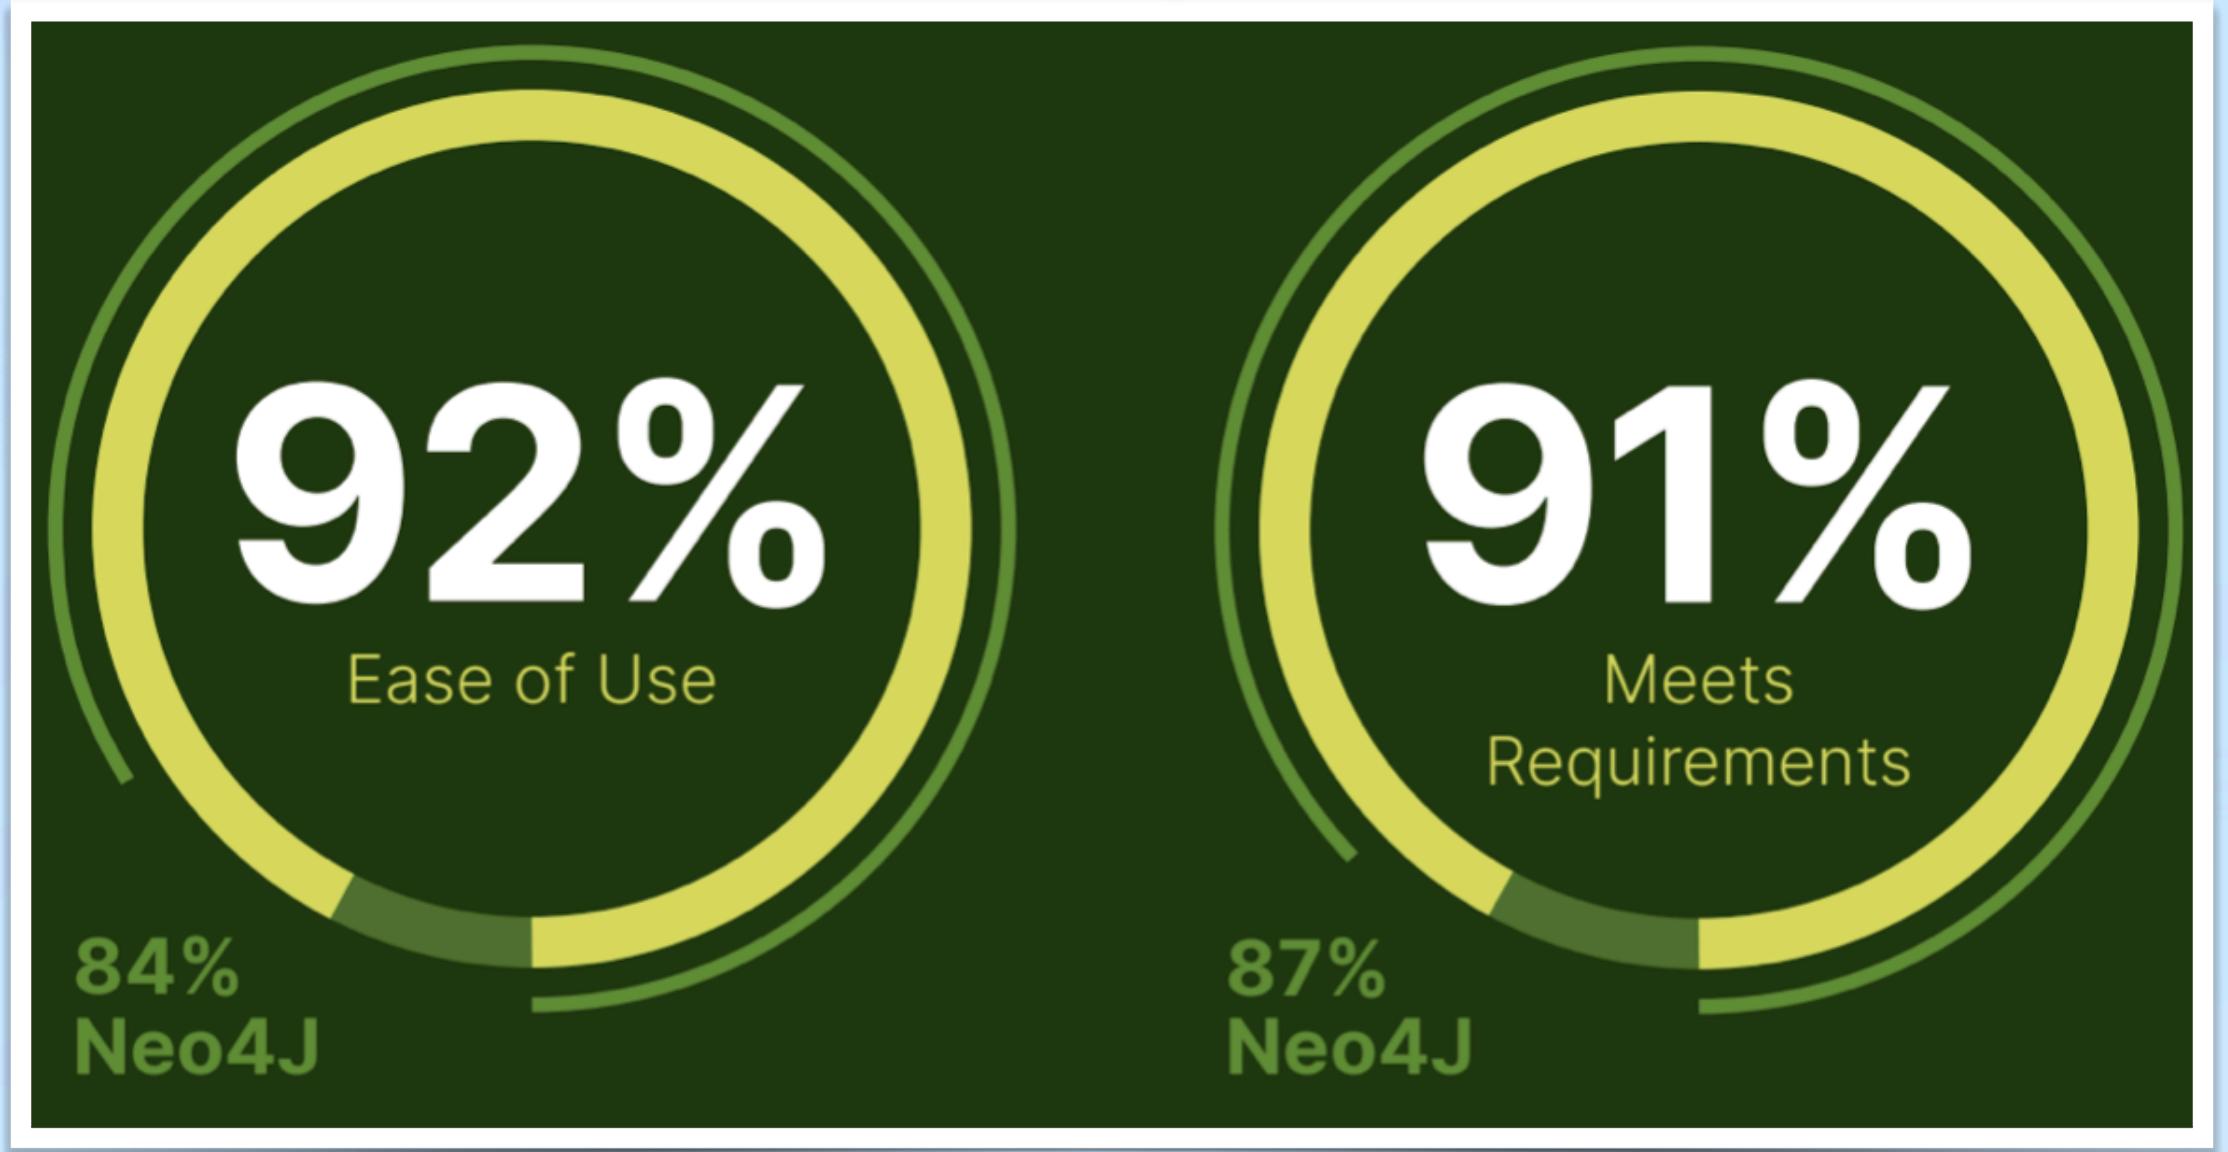
Review source: Seller invite

Incentivized Review



Oct 27, 2023

"Great DB for real world data"



Meets Requirements	8.8 <small>⑧ 99</small>	8.5 <small>⑧ 72</small>
Ease of Use	8.9 <small>⑧ 100</small>	8.2 <small>⑧ 74</small>
Ease of Setup	8.7 <small>⑧ 67</small>	8.8 <small>⑧ 30</small>
Ease of Admin	8.7 <small>⑧ 54</small>	8.1 <small>⑧ 27</small>
Quality of Support	8.9 <small>⑧ 87</small>	8.5 <small>⑧ 63</small>



License

From **Apache 2.0** license to **BSL 1.1** to **Apache 2.0** again after 4 years

Current license allows full usage of the ArangoDB source code for any purpose except for providing a managed service of ArangoDB.

ArangoDB Community Edition - commercial use up to 100GB limit on dataset size within a single cluster

Implementation

Docker

```
services:  
  arangodb:  
    image: arangodb  
    container_name: arangodb  
    environment:  
      - ARANGO_ROOT_PASSWORD=test123  
  ports:  
    - "8529:8529"  
  volumes:  
    - arangodb_data:/var/lib/arangodb3  
    - ./data:/data  
  
volumes:  
  arangodb_data:  
    driver: local
```

Entering database

```
/ # arangosh --server.username root --server.password test123

arangosh (ArangoDB 3.12.2 [linux] 64bit, using jemalloc, build
refs/tags/v3.12.2 154c4c5632c, VPack 0.2.1, RocksDB 7.2.0, ICU
64.2, V8 12.1.165, OpenSSL 3.3.1 4 Jun 2024, build-id:
0f674dfac5247637960930fec4f02acaf5016b8)
Copyright (c) ArangoDB GmbH
This executable uses the GNU C library (glibc), which is
licensed under the GNU Lesser General Public License (LGPL),
see https://www.gnu.org/copyleft/lesser.html and
https://www.gnu.org/licenses/gpl.html

Command-line history will be persisted when the shell is
exited. You can use `--console.history false` to turn this off
Connected to ArangoDB 'http+tcp://127.0.0.1:8529, version:
3.12.2 [SINGLE, server], database: '_system', username: 'root'

Type 'tutorial' for a tutorial or 'help' to see common examples
127.0.0.1:8529@_system>
```

Javascript API

```
127.0.0.1:8529@_system> db._databases();
[
  "_system",
  "wlgs"
]

127.0.0.1:8529@_system> db._useDatabase("wlgs");
true

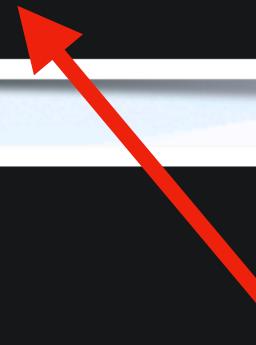
127.0.0.1:8529@wlgs> db._collections();
[
  [ArangoCollection 41321052, "_analyzers" (type document, status loaded)],
  [ArangoCollection 41321057, "_appbundles" (type document, status loaded)],
  [ArangoCollection 41321056, "_apps" (type document, status loaded)],
  [ArangoCollection 41321053, "_aqlfunctions" (type document, status loaded)],
  [ArangoCollection 41321058, "_frontend" (type document, status loaded)],
  [ArangoCollection 41321051, "_graphs" (type document, status loaded)],
  [ArangoCollection 41321055, "_jobs" (type document, status loaded)],
  [ArangoCollection 41321054, "_queues" (type document, status loaded)],
  [ArangoCollection 41321092, "edges" (type edge, status loaded)],
  [ArangoCollection 41321087, "nodes" (type document, status loaded)]
]

127.0.0.1:8529@wlgs> ('b' + 'a' + + 'a' + 'a').toLowerCase()
banana
```

DB initialization

```
● ● ●  
db._dropDatabase("wlgs");  
db._createDatabase("wlgs");  
db._useDatabase("wlgs");  
db._createDocumentCollection("nodes");  
db._createEdgeCollection("edges");  
var graph_module = require("@arangodb/general-graph");  
var graph = graph_module._create("graph");  
graph;  
  
graph._addVertexCollection("nodes");  
var rel = graph_module._relation("edges", ["nodes"], ["nodes"]);  
graph._extendEdgeDefinitions(rel);
```

```
● ● ●  
#!/bin/sh  
arangosh --server.username root --server.password test123 --javascript.execute /data/init.js  
  
arangoimport --collection nodes \  
  --file "./data/nodes.csv" \  
  --type csv \  
  --server.database wlgs \  
  --server.endpoint http://localhost:8529 \  
  --server.username root \  
  --server.password test123 \  
  --on-duplicate ignore  
  
arangoimport --collection edges \  
  --file "./data/edges.csv" \  
  --type csv \  
  --server.database wlgs \  
  --server.endpoint http://localhost:8529 \  
  --server.username root \  
  --server.password test123 \  
  --on-duplicate ignore
```



Import time

```
● ● ●  
> time python ./preprocess_csv.py  
⌚ Fixing and encoding ./data/popularity_iw.csv... ✓ 1.12s.  
⌚ Fixing and encoding ./data/taxonomy_iw.csv... ✓ 10.81s.  
⌚ Processing ./data/popularity_iw.csv... ✓ 1.13s.  
⌚ Processing ./data/taxonomy_iw.csv... ✓ 16.67s.  
28.59s user 0.87s system 98% cpu 29.772 total
```

```
● ● ●  
# time ./data/init.sh  
. . .  
created: 5483248  
warnings/errors: 288361  
updated/replaced: 0  
ignored: 0  
lines read: 5771611  
real 7m 15.72s  
user 0m 4.29s  
sys 0m 2.41s
```

Python driver

```
● ● ●

from arango import ArangoClient

client = ArangoClient(hosts="http://localhost:8529")
db = client.db("wlgs", username="root", password="test123")

aql_query = """
UPDATE @node_key WITH { popularity: @new_popularity } IN nodes
RETURN OLD
"""

cursor = db.aql.execute(
    aql_query,
    bind_vars={'node_key': node_key,
               'new_popularity': str(new_popularity)}
)
```

```
● ● ●

aql_query = """
FOR v IN nodes
    LET incoming = (
        FOR vertex IN 1..1 INBOUND v GRAPH 'graph'
            RETURN vertex
    )
    FILTER LENGTH(incoming) == 0
    RETURN {
        key: v._key,
        title: v.title,
        popularity: v.popularity
    }
"""
cursor = db.aql.execute(aql_query)
results = [doc for doc in cursor]
```

Cool stuff

K-paths

Task 14 was a breeze

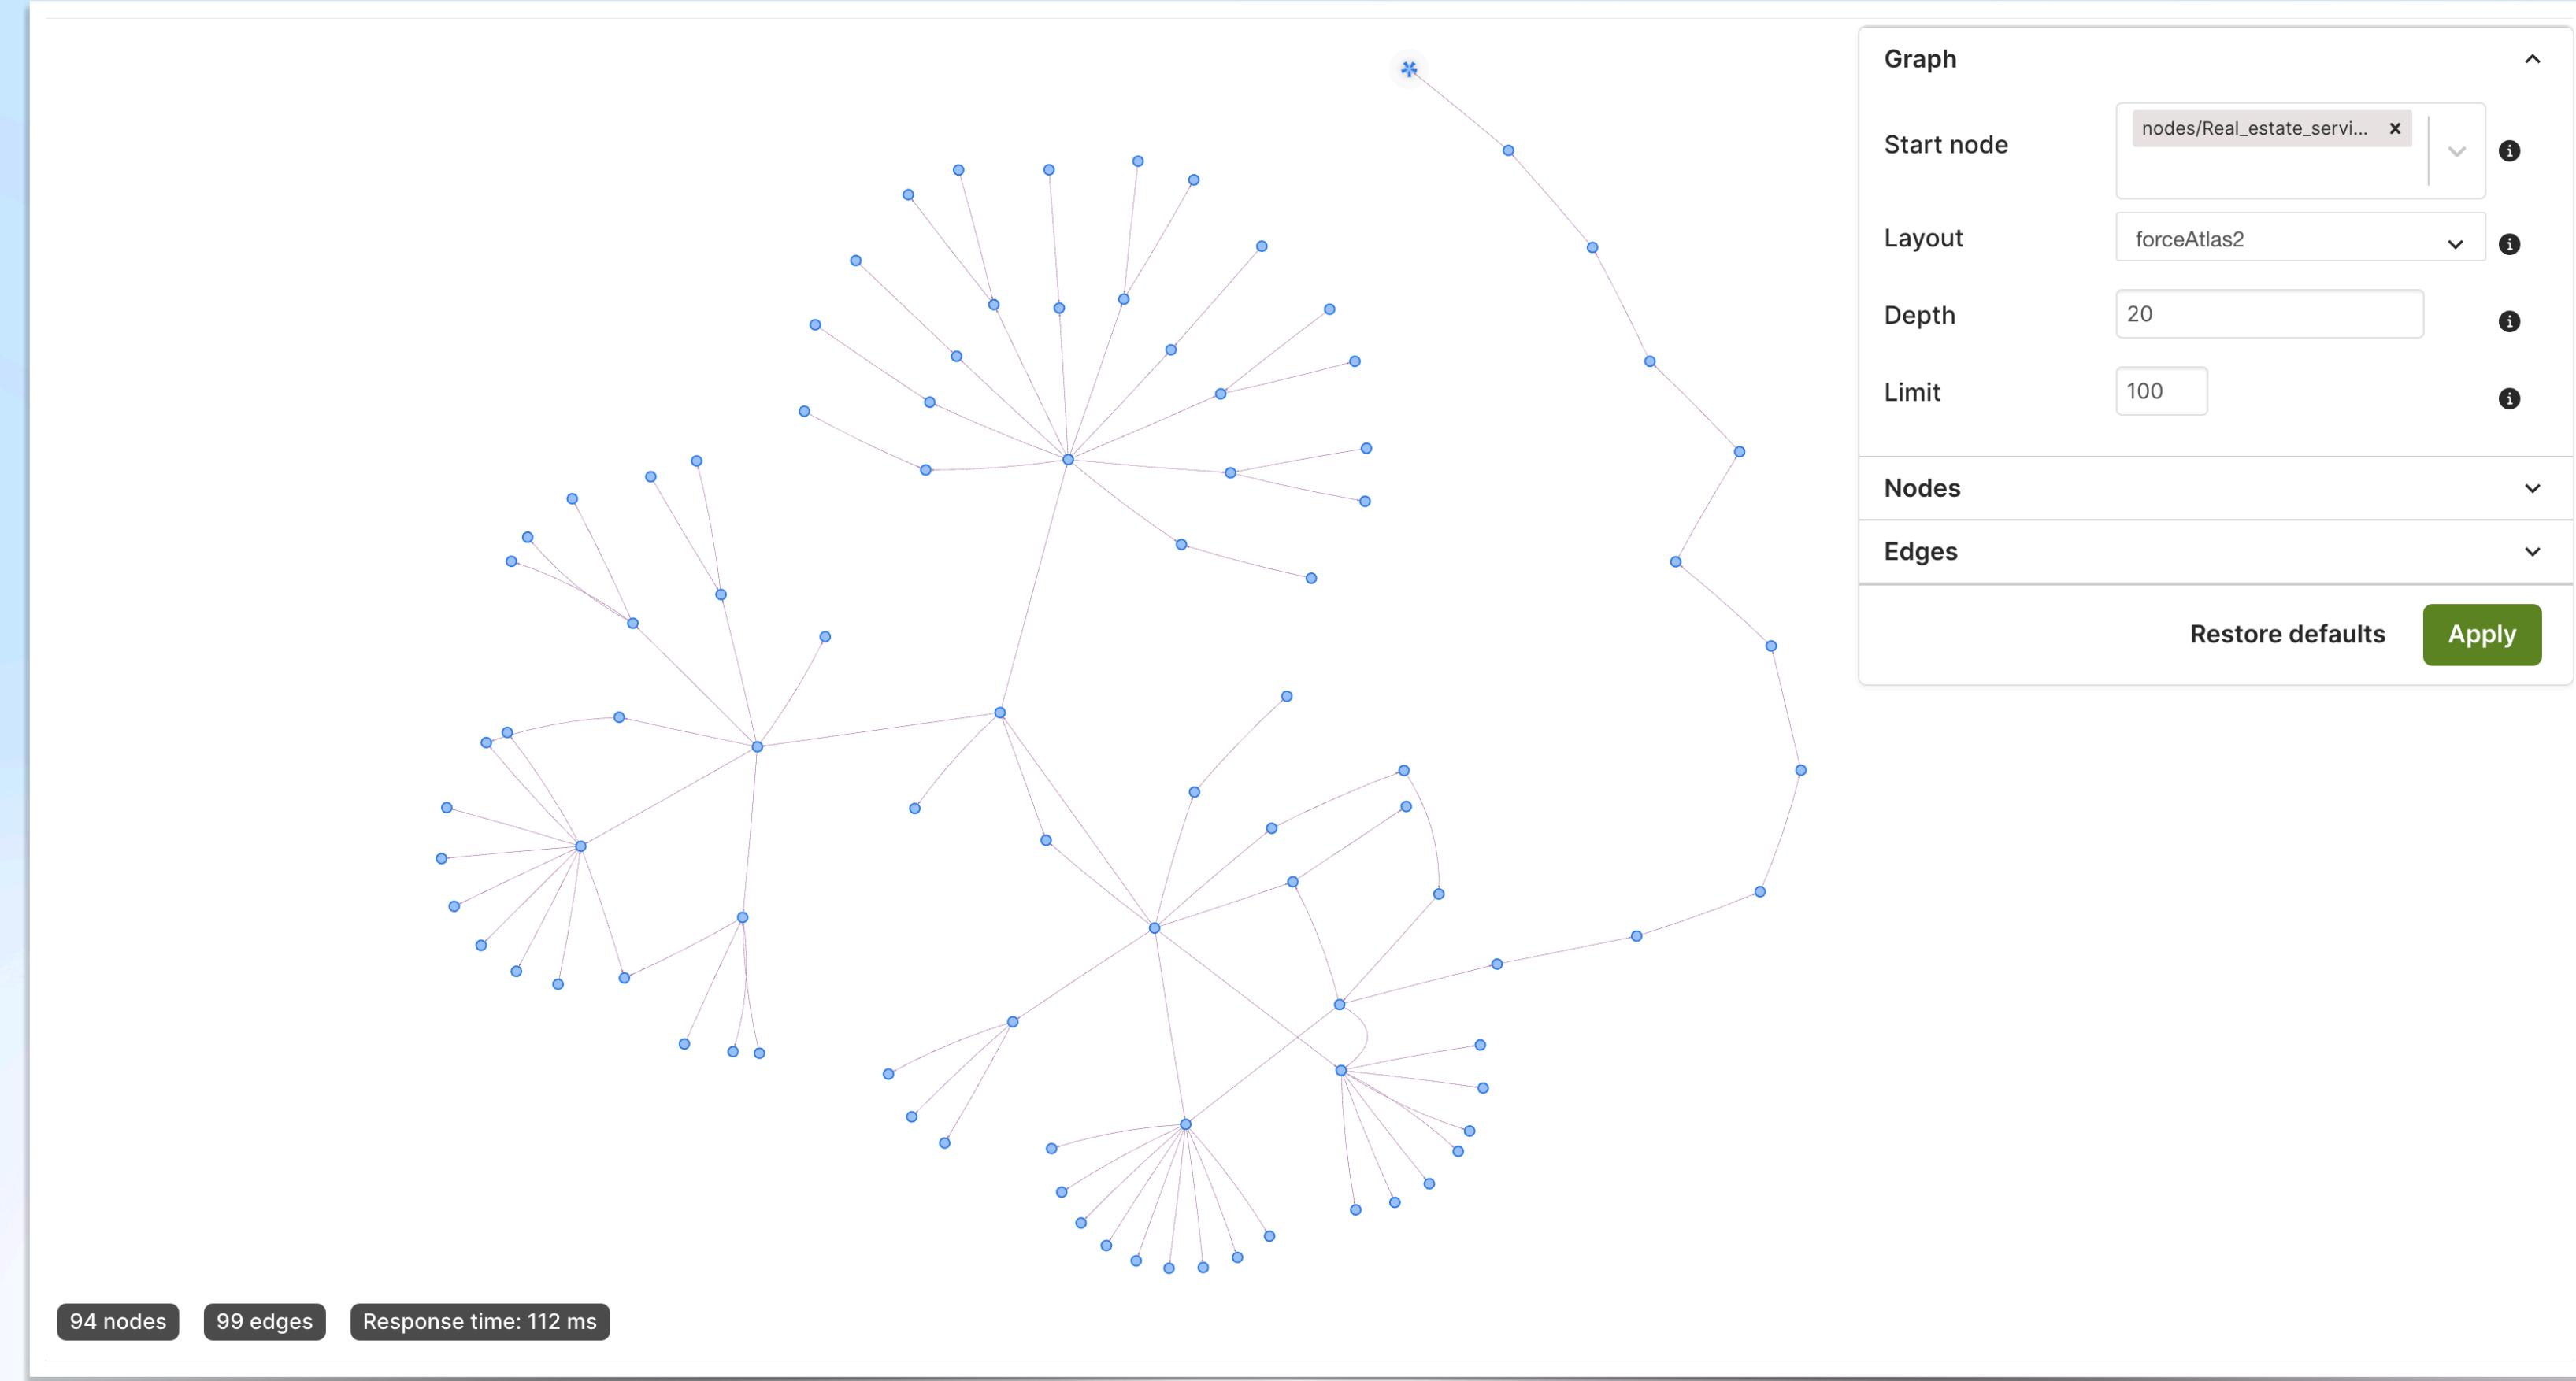
```
FOR path
    IN MIN..MAX OUTBOUND|INBOUND|ANY K_PATHS
    startVertex TO targetVertex
    GRAPH graphName
    [OPTIONS options]
```

.getExtra()

```
127.0.0.1:8529@wlgs> db._query(`  
...>   LET neighborhood = (  
...>     FOR v, e, p IN 0..@radius ANY @start_vertex  
...>       GRAPH 'graph'  
...>       RETURN DISTINCT {  
...>         key: v._key,  
...>         popularity: TO_NUMBER(v.popularity)  
...>       }  
...>     )  
...>   RETURN {  
...>     total_popularity: SUM(neighborhood[*].popularity),  
...>     node_count: LENGTH(neighborhood)  
...>   }  
...> `,  
...> {radius: 5, start_vertex: 'nodes/Microeconomics' },  
...> {},  
...> { fullCount: true }  
...> ).getExtra();
```

```
{  
  "warnings" : [ ],  
  "stats" : {  
    "writesExecuted" : 0,  
    "writesIgnored" : 0,  
    "documentLookups" : 0,  
    "seeks" : 0,  
    "scannedFull" : 0,  
    "scannedIndex" : 5934909,  
    "cursorsCreated" : 2,  
    "cursorsRearmed" : 251016,  
    "cacheHits" : 171576,  
    "cacheMisses" : 79442,  
    "filtered" : 125560,  
    "httpRequests" : 0,  
    "fullCount" : 1,  
    "executionTime" : 13.531461615000808,  
    "peakMemoryUsage" : 222593024,  
    "intermediateCommits" : 0  
  }  
}
```

Visualization*



*ArangoDB Community Edition feature

Not so cool stuff

Write operation locks collection for read

Design limitations

The following design limitations are known for AQL queries:

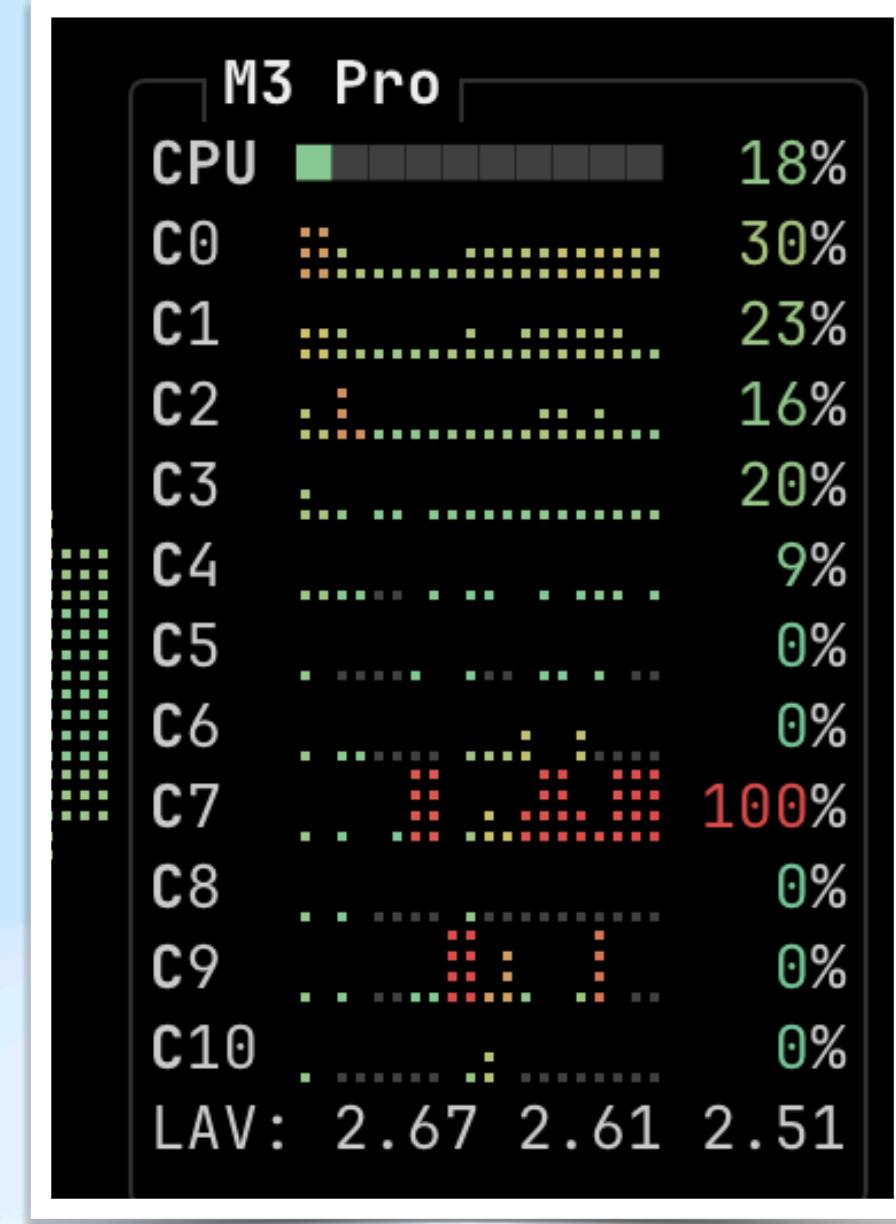
- Up to v3.12.0, subqueries used inside expressions are pulled out of these expressions and executed beforehand. That means that subqueries do not participate in lazy evaluation of operands, for example in the [ternary operator](#) or when used as sub-expressions that are combined with logical `AND` or `OR`.

From v3.12.1 onward, short-circuiting is applied.

Also see [evaluation of subqueries](#).

- It is not possible to use a collection in a read operation after it was used for a write operation in the same AQL query.
- In the cluster, all collections that are accessed **dynamically** by [traversals working with collection sets](#) (instead of named graphs) must be stated in the query's initial [WITH statement](#). To make the `WITH` statement required in single server as well (e.g. for testing a migration to cluster), please start the server with the option `--query.require-with`.

Single threaded queries



```
✓ Container arangodb Created 0.0s
Attaching to arangodb
arangodb | 2025-01-22T15:32:14.157980Z [1-2] INFO [e52b0] {general} ArangoDB 3.12.2 [linux] 64bit, using jemalloc, build refs/tags/v3.12.2
154c4c5632c, VPack 0.2.1, RocksDB 7.2.0, ICU 64.2, V8 12.1.165, OpenSSL 3.3.1 4 Jun 2024, build-id: 4a5d39cb634736e1e6dc8c4f25b92a395087b915
arangodb | 2025-01-22T15:32:14.157989Z [1-2] INFO [11111] {general} This executable uses the GNU C library (glibc), which is licensed under the
GNU Lesser General Public License (LGPL), see https://www.gnu.org/copyleft/lesser.html and https://www.gnu.org/licenses/gpl.html
arangodb | 2025-01-22T15:32:14.158244Z [1-2] INFO [a1c60] {syscall} file-descriptors (nofiles) hard limit is 1048576, soft limit is 1048576
arangodb | 2025-01-22T15:32:14.158340Z [1-2] INFO [75ddc] {general} detected operating system: Linux version 6.11.5-orbstack-00280-g96d99c92a42b
(arbstack@builder) (ClangBuiltLinux clang version 18.1.8 (https://github.com/llvm/llvm-project.git 3b5b5c1ec4a3095ab096dd780e84d7ab81f3d7ff>),
ClangBuiltLinux LLD 18.1.8) #51 SMP Sun Nov 3 08:07:37 UTC 2024
arangodb | 2025-01-22T15:32:14.158370Z [1-2] INFO [25362] {memory} Available physical memory: 6275833856 bytes, available cores: 11
```

Key limitations

- It must consist of the letters A to Z (lower- and uppercase), the digits 0 to 9, or any of the following punctuation characters: _ - . @ () + , = ; \$! * ' % :

 Avoid using : in document keys as this can conflict with the requirements of [EnterpriseGraphs](#), [SmartGraphs](#), and [SmartJoins](#) which use the colon character as a separator in keys.

- Any other characters, especially multi-byte UTF-8 sequences, whitespace, or punctuation characters not listed above cannot be used inside key values.

```
SAFE_CHARACTERS = "_-.@( )+=;:$!*'%:"  
key = urlencode("ÓΣεάσθ&żć®łń©kvl"'^óΔŻłńĶ", safe=SAFE_CHARACTERS)  
key = key.replace('~', '%7E')  
key = key.replace(':', '%3A')
```

Memory

```
/app # ./dbcli 14 Microeconomics Service_companies

Error finding paths: [HTTP 500][ERR 32] AQL: global memory limit
exceeded [node #11: SubqueryStartNode] [node #4: CalculationNode]
[node #5: EnumerateListNode] [node #6: CalculationNode] [node #12:
SubqueryEndNode] [node #9: CalculationNode] [node #9:
CalculationNode] [node #10: ReturnNode] (while executing)
```

```
FOR path
IN 1..999 OUTBOUND K_PATHS
@start_vertex T0 @end_vertex
GRAPH 'graph'
OPTIONS {
    uniqueVertices: 'path',
}
RETURN {
    vertices: (
        FOR vertex IN path.vertices
        RETURN {
            key: vertex._key,
            title: vertex.title
        }
    ),
    edges: LENGTH(path.edges),
    weight: path.weight
}
```

Thank you